

Building your First Apache Kafka Application



Bogdan Sucaciu

SOFTWARE ENGINEER @ AXUAL

@BSucaciu

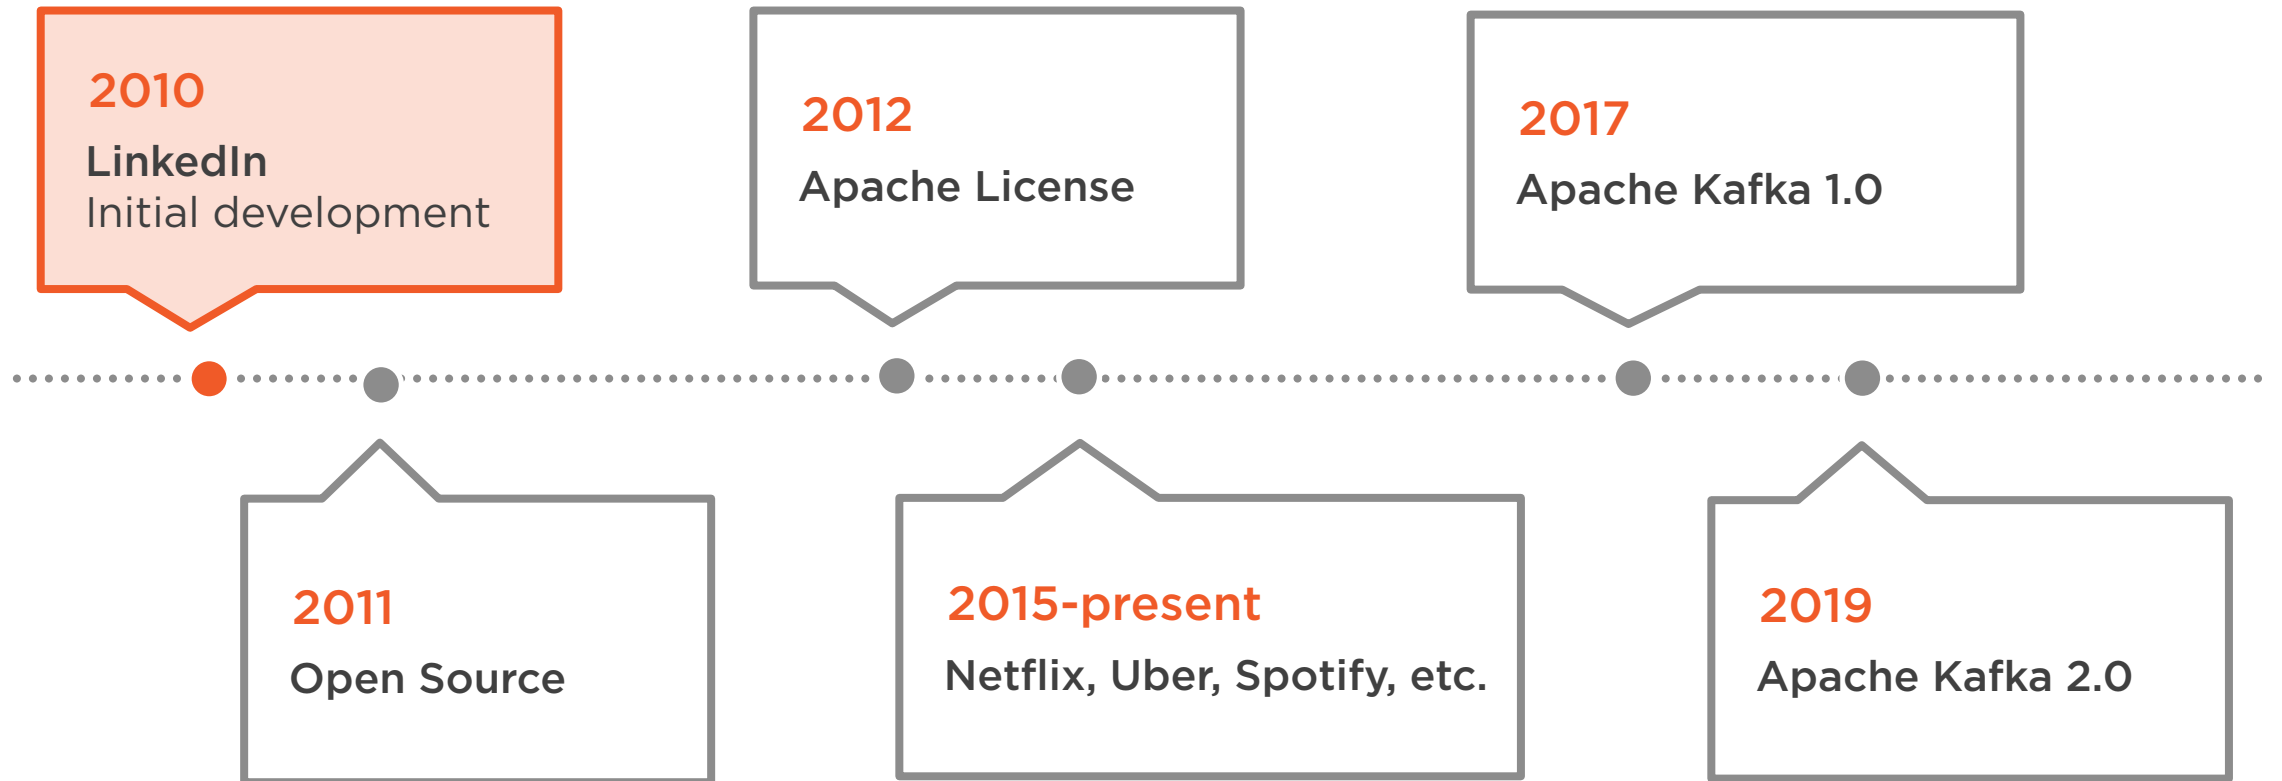


Why Kafka?

Open source



Kafka History



Why Kafka?

Open source

Java



Classes




Scala




Java


Compiler



Bytecode



Bytecode



Why Kafka?

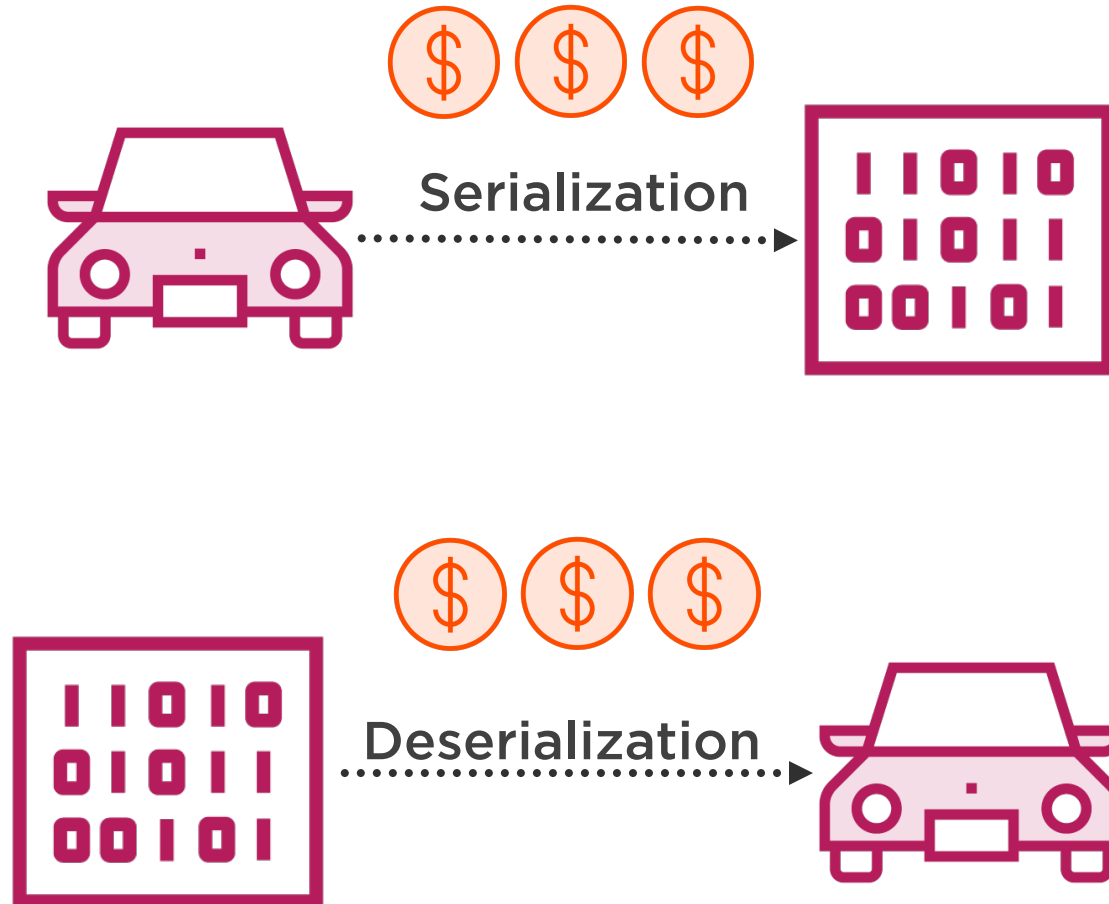
Open source

Java

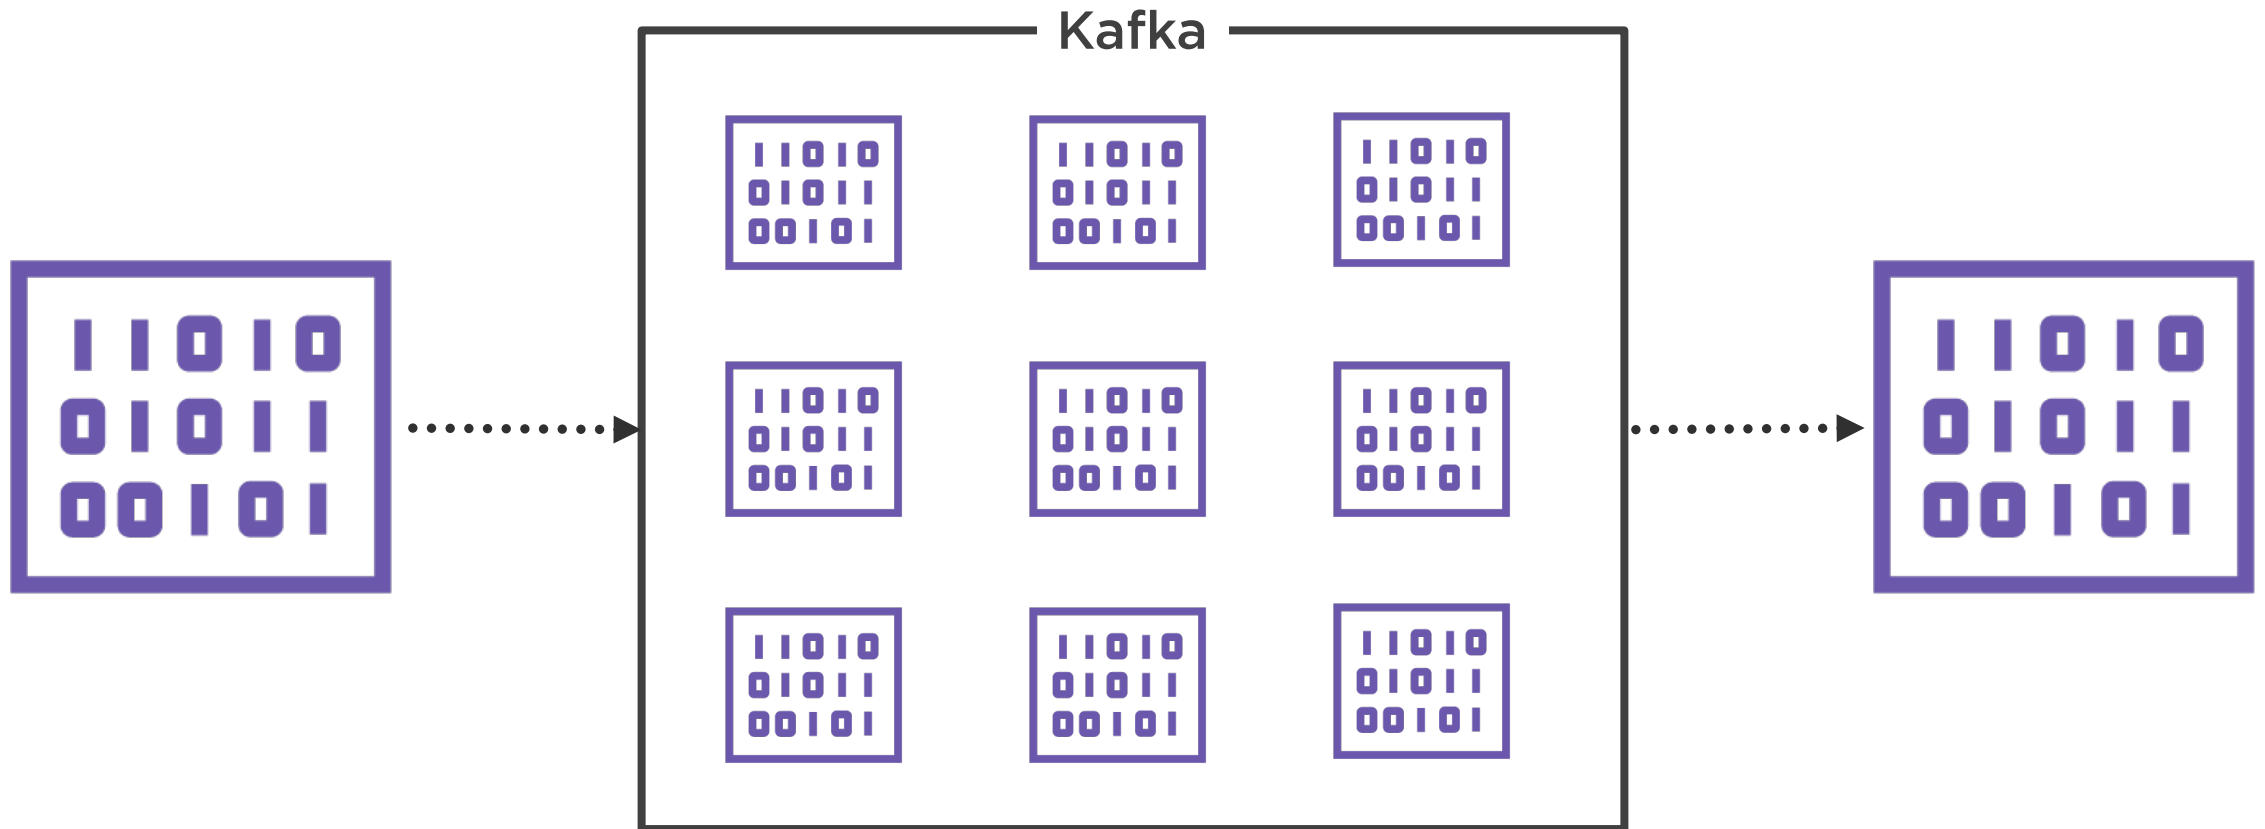
High throughput




No Serialization / Deserialization




No Serialization / Deserialization



Zero Copy



Zero Copy



Why Kafka?

Open source

Java

High throughput


More than a messaging system




Distributed Streaming Platform



Messaging system
Publish / Subscribe
pattern



Distributed storage
Store data in a fault-
tolerant, durable way



Data processing
Process events as they
occur


Kafka Components

Broker




Zookeeper




Broker



Kafka Cluster



Zookeeper



Demo



Setup a 2 node Kafka cluster:

- 2x Broker
- 1x Zookeeper



Record

Key

Value

Timestamp

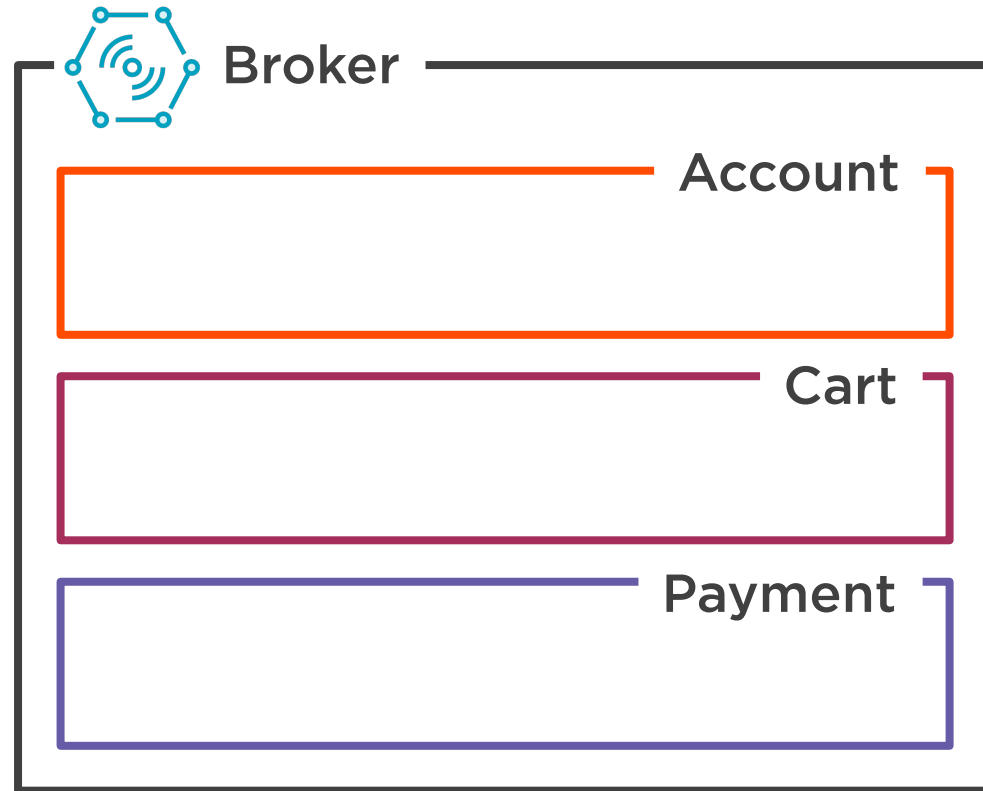
Anything

Anything

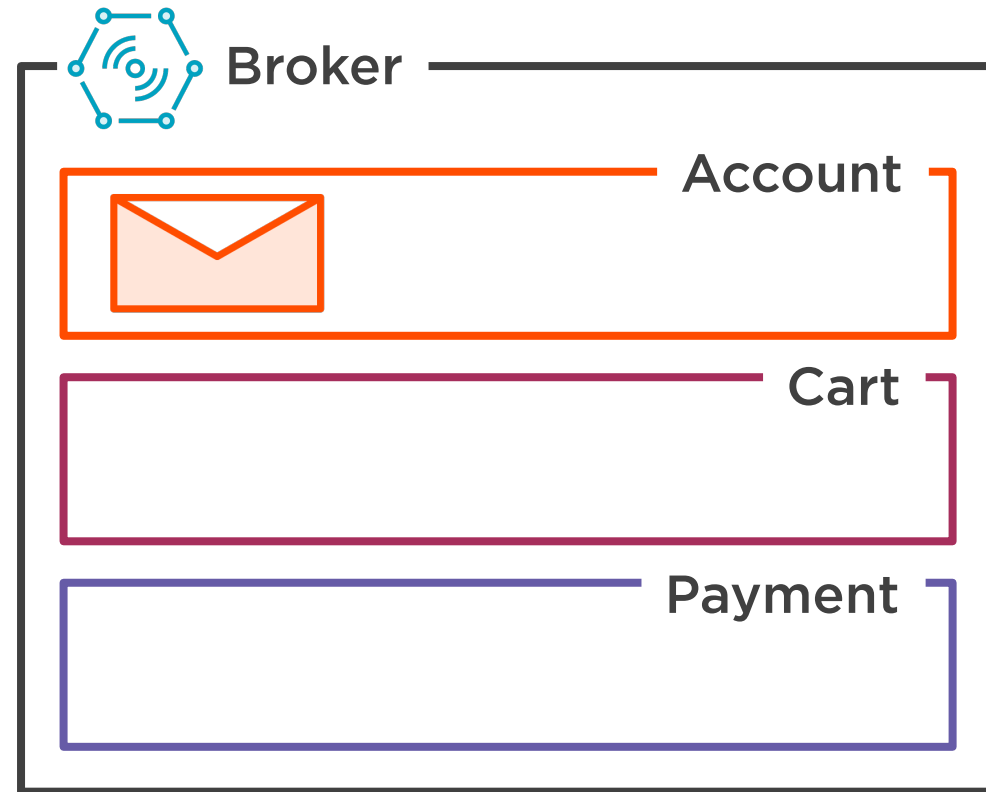
Timestamp format



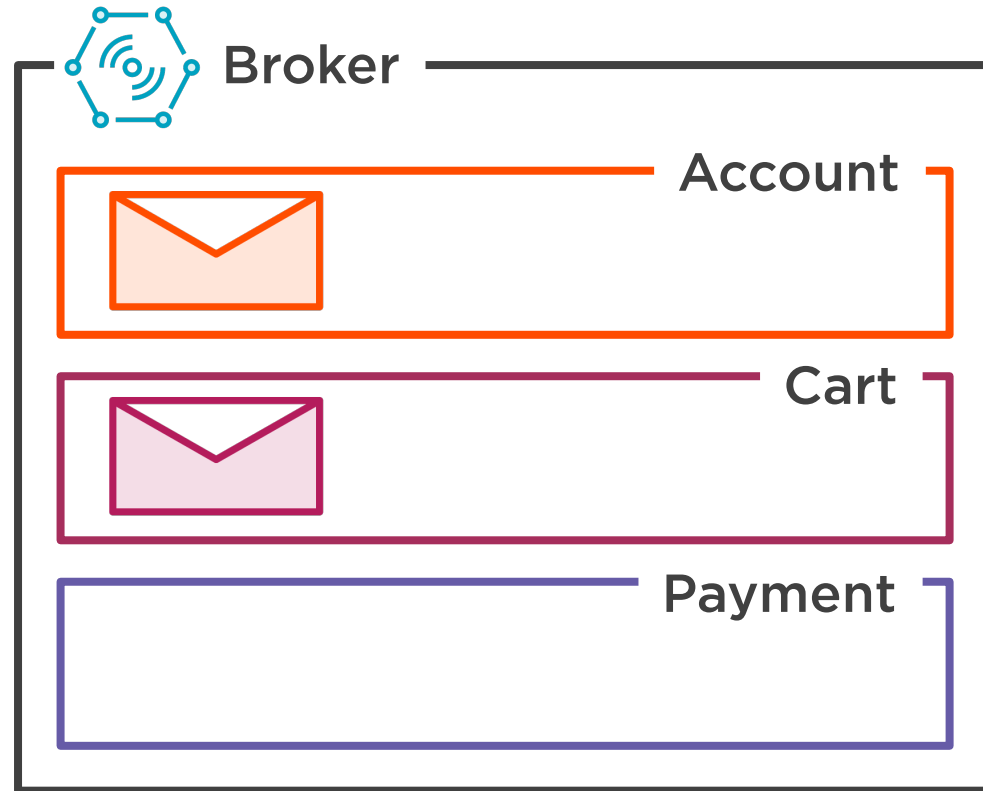
Topics



Topics



Topics



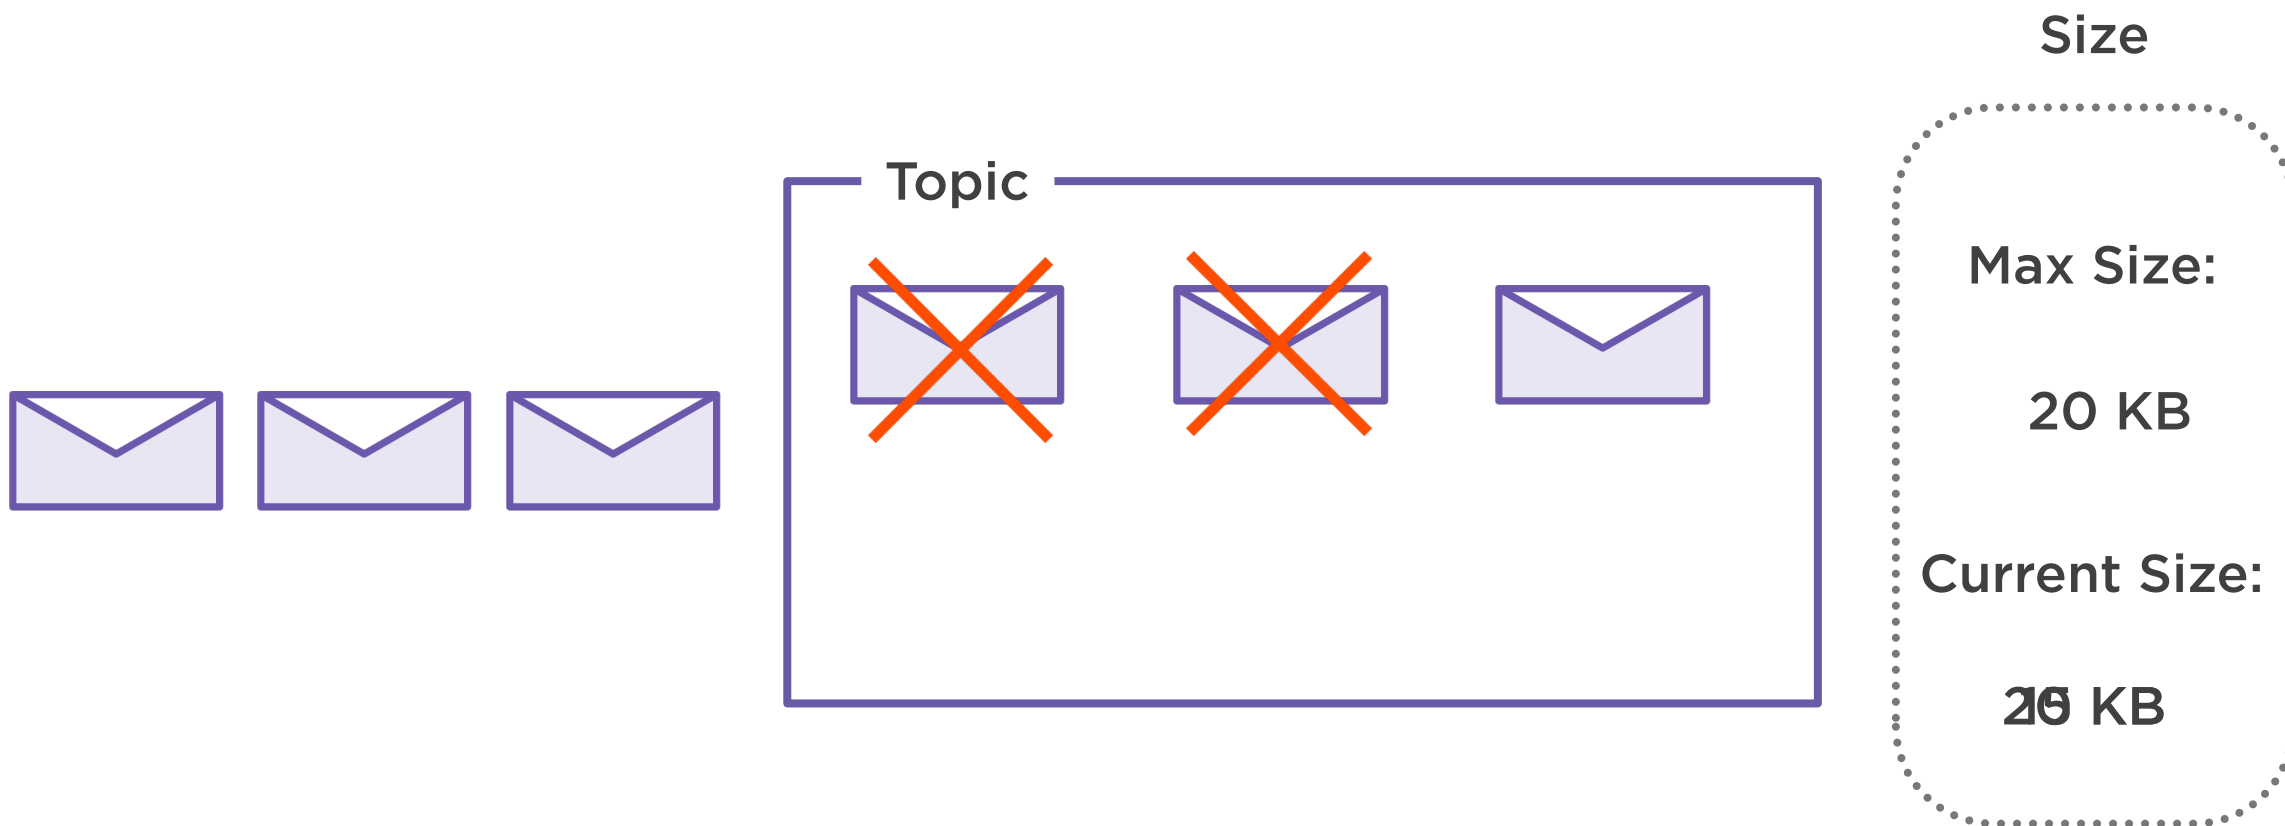
Topics

Delete


Compact




Delete Topics



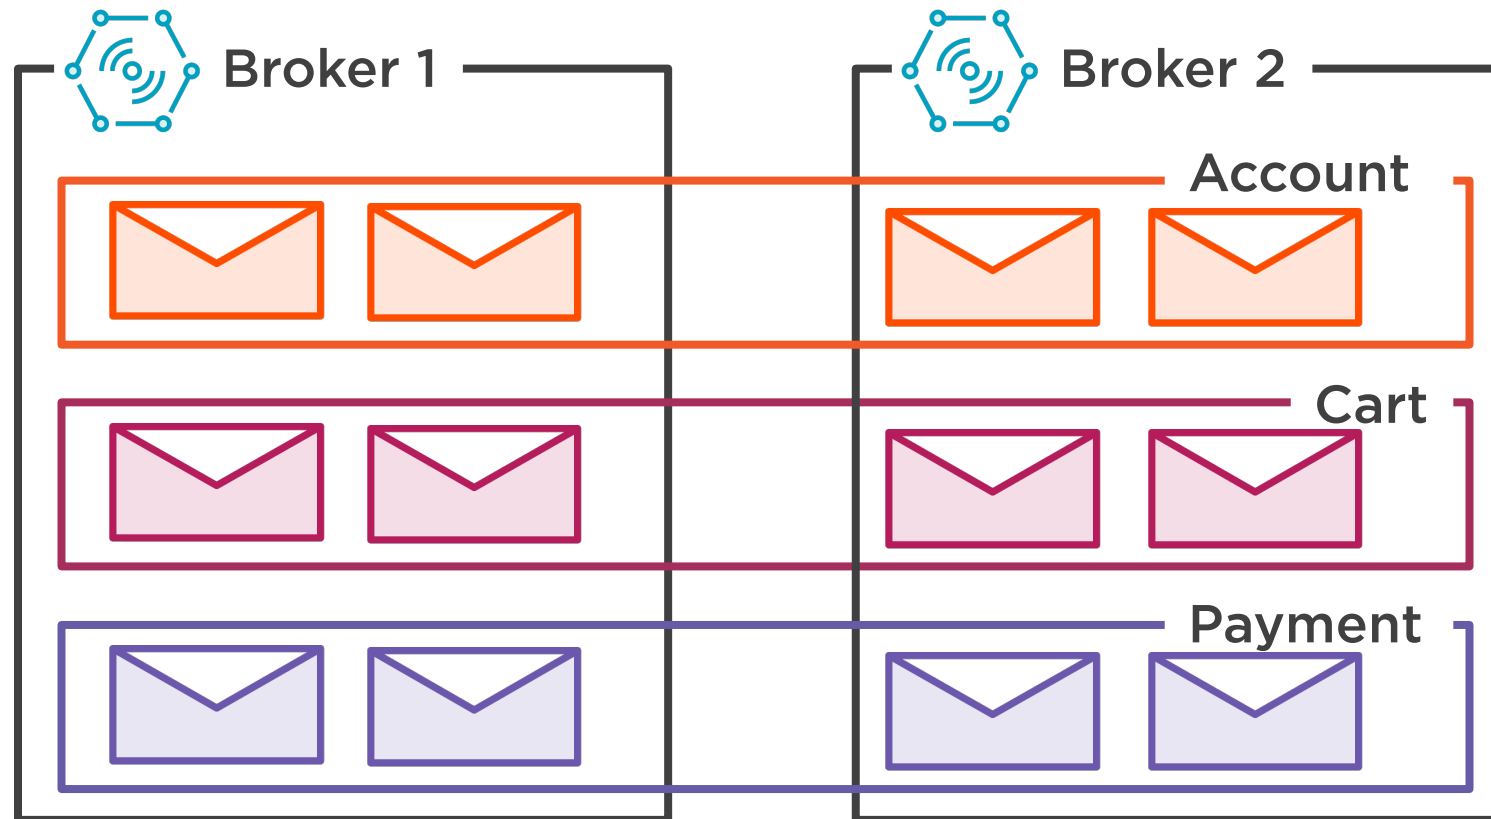
Delete Topics





Compaction Topics




Topics



Partitions




Replicated Partitions




Know Your Data !



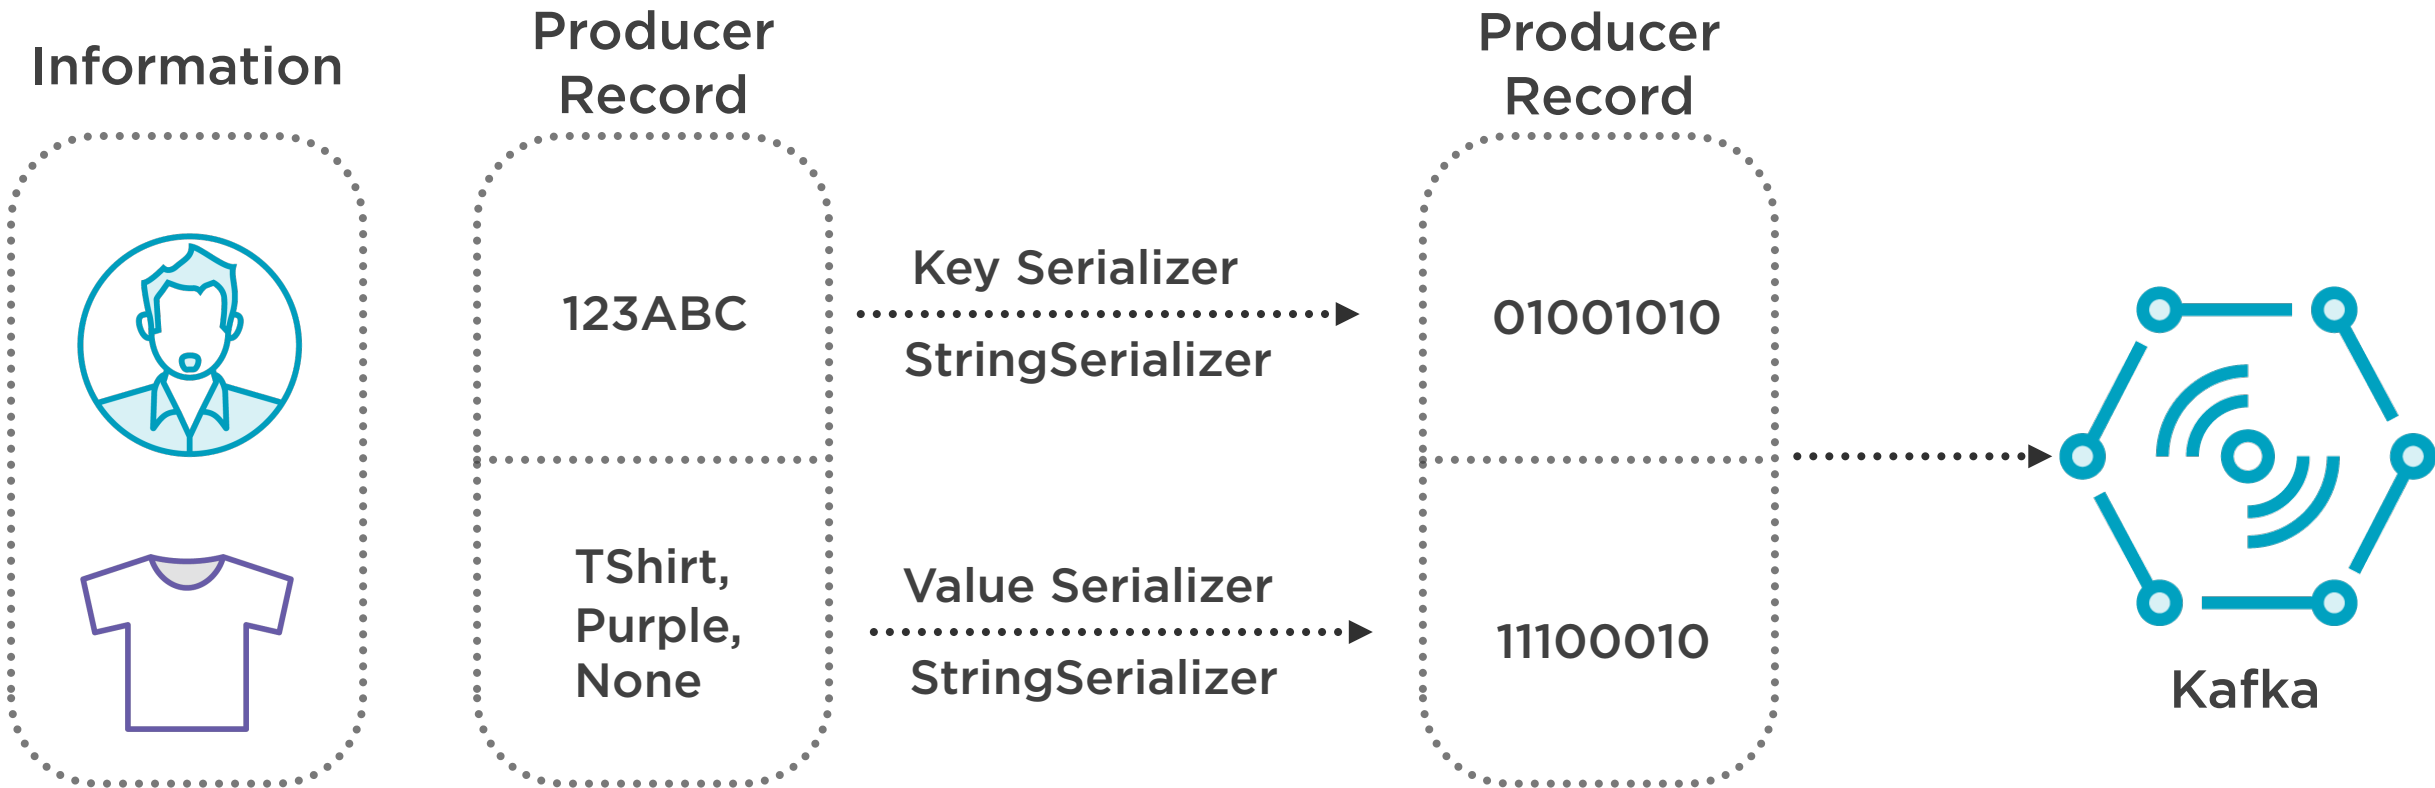
Producer



Producer



Producer



Demo



Basic producer application

Prerequisites:

- Java 8+
- Maven



Consumer

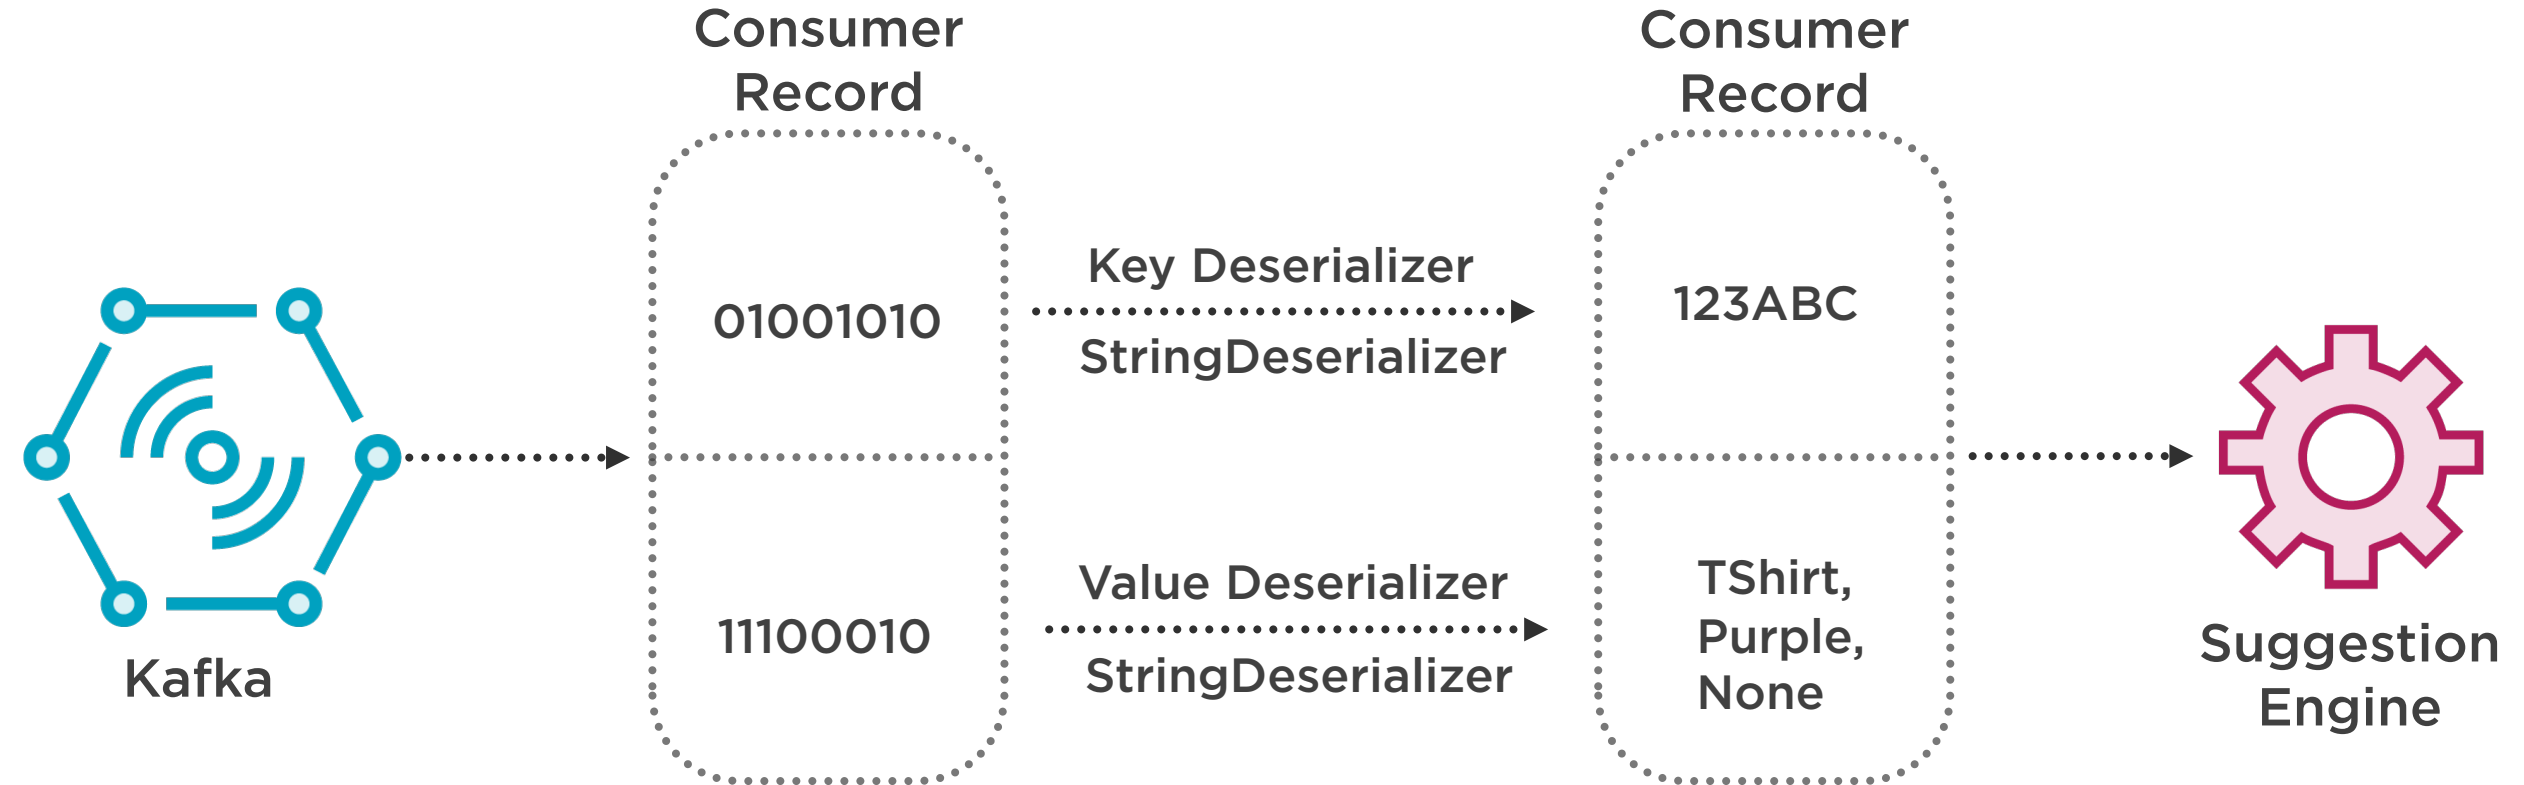


Kafka



Consumer

Consumer



Demo



Basic consumer application



Summary



Why Kafka?

Kafka Architecture

Topics

Kafka Clients

- Producer
- Consumer

